

PARTNERSHIP FOR ADVANCED COMPUTING IN EUROPE

Introduction to OpenMP





Introduction to OpenMP





Outline

- What is OpenMP?
- Timeline
- Main Terminology
- OpenMP Programming Model
- Main Components
- Parallel Construct
- Work-sharing Constructs
 - sections, single, workshare
- Data Clauses
 - default, shared, private, firstprivate, lastprivate, threadprivate, copyin





What is OpenMP?

OpenMP (*Open specifications for Multi Processing*)

- is an API for shared-memory parallel computing;
- is an open standard for portable and scalable parallel programming;
- is flexible and easy to implement;
- is a specification for a set of compiler directives, library routines, and environment variables;
- is designed for C, C++ and Fortran.





Timeline



- OpenMP 4.0 Release Candidate 1 was released in November 2012.
- <http://openmp.org/>





Main Terminology



1. OpenMP thread: a **lightweight** process
2. thread team: a set of threads which co-operate on a task
3. master thread: the thread which co-ordinates the team
4. thread-safety: correctly executed by multiple threads
5. OpenMP directive: line of code with meaning only to certain compilers
6. construct: an OpenMP executable directive
7. clause: controls the scoping of variables during the execution





OpenMP Programming Model

OpenMP is designed for multi-processor/core UMA or NUMA shared memory systems.



Execution Model:

- Thread-based Parallelism
- Compiler Directive Based
- Explicit Parallelism
- Fork-Join Model



- Dynamic Threads
- Nested Parallelism




Memory Model:

- All threads have access to the shared memory.
- Threads can share data with other threads, but also have private data.
- Threads sometimes **synchronise** against **data race**.
- Threads cache their data; Use **OpenMP** **flush**

Thread 1

Thread 2

Thread 3





Main Components

- Compiler Directives and Clauses: appear as comments, executed when the appropriate OpenMP flag is specified
 - Parallel construct
 - Work-sharing constructs
 - Synchronization constructs
 - Data Attribute clauses

C/C++: ***#pragma omp directive-name [clause[clause]...]***

Fortran free form: ***!\$omp directive-name [clause[clause]...]***

Fortran fixed form: ***!\$omp | c\$omp | *\$omp directive-name [clause[clause]...]***





Compiling:

	Compiler	Flag
Intel	icc (C)	-openmp
	icpc (C++)	
	ifort (Fortran)	
GNU	gcc (C)	-fopenmp
	g++ (C++)	
	g77/gfortran (Fortran)	
PGI	pgcc (C)	-mp
	pgCC (C++)	
	pg77/pgfortran (Fortran)	

See: <http://openmp.org/wp/openmp-compilers/> for the full list.





- Runtime Functions: for managing the parallel program
 - `omp_set_num_threads(n)` - set the desired number of threads
 - `omp_get_num_threads()` - returns the current number of threads
 - `omp_get_thread_num()` - returns the id of this thread
 - `omp_in_parallel()` – returns .true. if inside parallel region and more.

For C/C++: Add `#include <omp.h>`

For Fortran: Add use `omp_lib`

- Environment Variables: for controlling the execution of parallel program at run-time.

- csh/tcsh: `setenv OMP_NUM_THREADS n`
- ksh/sh/bash: `export OMP_NUM_THREADS=n`
- and more.





Parallel Construct

- The fundamental construct in OpenMP.
 - Every thread executes the same statements which are inside the parallel region simultaneously.
 - At the end of the parallel region there is an implicit barrier for synchronization
- c/c++:

<code>#pragma omp parallel [clauses]</code>	<code>!\$omp parallel [clauses]</code>
<code>{</code>	<code>...</code>
<code>...</code>	<code>!\$omp end</code>
<code>}</code>	<code>parallel</code>



- Create a 4-thread parallel region

- Each thread with tid from 0 to 3 calls `foo(tid, A)`

- Threads wait for all threads to finish before proceeding

```

double A[1000];
omp_set_num_threads(4);
#pragma omp parallel
{
  int tid=omp_get_thread_num();
  foo(tid,A);
}
printf( "All Done\n" );
  
```

double A[1000];

omp_set_num_threads(4);

foo(0,A); foo(1,A); foo(2,A); foo(3,A);

printf("All Done\n");





Hello World Example:

```
C:
#include<omp.h>
#include<stdio.h>

int main(){
#pragma omp parallel

printf("Hello from thread %d out
of %d\n", omp_get_thread_num(),
omp_get_num_threads());
}
```

```
Fortran:
program hello
use omp_lib

implicit none
!$omp parallel

PRINT*, 'Hello from
thread',omp_get_thread_num(),'out
of',omp_get_num_threads()

!$omp end parallel

end program hello
```





Compile: (Intel)

```
>icc -openmp hello.c -o a.out  
>ifort -openmp hello.f90 -o a.out
```

Execute:

```
>export OMP_NUM_THREADS=4  
>./a.out  
Hello from thread 0 out of 4  
Hello from thread 3 out of 4  
Hello from thread 1 out of 4  
Hello from thread 2 out of 4
```





- **Dynamic threads:**

- The number of threads used in a parallel region can vary from one parallel region to another.
- `omp_set_dynamic()`, `OMP_DYNAMIC`
- `omp_get_dynamic()`

- **Nested parallel regions:**

- If a parallel directive is encountered within another parallel directive, a new team of threads will be created.
- `omp_set_nested()`, `OMP_NESTED`
- `omp_get_nested()`





- **If Clause:**

- Used to make the parallel region directive itself conditional.
- Only execute in parallel if expression is true.

c/C++:

```
#pragma omp parallel if(n>100)
{
    ...
}
```

(Checks the size
of the data)

Fortran:

```
!$omp parallel if(n>100)
...
!$omp end parallel
```

- **nowait Clause:**

- allows threads that finish earlier to proceed without waiting

c/C++:

```
#pragma omp parallel nowait
{
    ...
}
```

```
!$omp parallel
...
!$omp end parallel
nowait
```





Data Clauses

- Used in conjunction with several directives to control the scoping of enclosed variables.
 - **default(*shared/private/none*)**: The default scope for all of the variables in the parallel region.
 - **shared(*/list*)**: Variable is shared by all threads in the team. All threads can read or write to that variable.
 - **private(*/list*)**: Each thread has a private copy of variable. It can only be read or written by its own thread.

C: #pragma omp parallel default(none), shared(n)

Fortran: !\$omp parallel default(none), shared(n)

C: #pragma omp parallel default(none), shared(n), private(tid)

Fortran: !\$omp parallel default(none), shared(n), private(tid)





- Most variables are shared by default
 - C/C++: File scope variables, static
 - Fortran: COMMON blocks, SAVE variables, MODULE variables
 - Both: dynamically allocated variables
- Variables declared in parallel region are always private
- How do we decide which variables should be shared and which private?
 - Loop indices - private
 - Loop temporaries - private
 - Read-only variables - shared
 - Main arrays - shared





Example:

```
C:
#include<omp.h>
#include<stdio.h>
int tid, nthreads;
int main(){
    #pragma omp parallel private(tid),
    shared(nthreads)
    {
        tid=omp_get_thread_num();
        nthreads=omp_get_num_threads();
        printf("Hello from thread %d out
        of %d\n", tid, nthreads);
    }
}
```

```
Fortran:
program hello
use omp_lib
implicit none
integer tid, nthreads

!$omp parallel private(tid),
shared(nthreads)
tid=omp_get_thread_num()
nthreads=omp_get_num_threads()
PRINT*, 'Hello from
thread',tid,'out of',nthreads
!$omp end parallel

end program hello
```





Some Additional Data Clauses:

- `firstprivate(list)`: Private copies of a variable are initialized from the original global object.
- `lastprivate(list)`: On exiting the parallel region, variable has the value that it would have had in the case of serial execution.
- `threadprivate(list)`: Used to make global file scope variables (C/C++) or common blocks (Fortran) local.
- `copyin(list)`: Copies the threadprivate variables from master thread to the team threads.
- `copyprivate` and `reduction` clauses will be described later.





Work-Sharing Constructs

- To distribute the execution of the associated region among threads in the team
- An implicit barrier at the end of the worksharing region, unless the nowait clause is added
- Work-sharing Constructs:
 - Loop
 - Sections
 - Single
 - Workshare





Sections Construct

- A non-iterative work-sharing construct.
- Specifies that the enclosed section(s) of code are to be executed by different threads.
- Each section is executed by one thread.

C/C++:


```
#pragma omp sections [clauses] nowait
{
    #pragma omp section
    ...
    #pragma omp section
    ...
}
```

Fortran:

```
!$omp sections [clauses]
!$omp section
...
!$omp section
...
!$omp end sections
[nowait]
```



```
#include <stdio.h>
#include <omp.h>
int main() {
    int tid;
    #pragma omp parallel private(tid)
    {
        tid=omp_get_thread_num();
        #pragma omp sections
        {
            #pragma omp section
            printf("Hello from thread %d \n", tid);
            #pragma omp section
            printf("Hello from thread %d \n", tid);
            #pragma omp section
            printf("Hello from thread %d \n", tid);
        }
    }
}
```



```
>export
OMP_NUM_THREADS=4
```

Hello from thread 0
Hello from thread 2
Hello from thread 3





Single Construct

- Specifies a block of code that is executed by only one of the threads in the team.
- May be useful when dealing with sections of code that are not thread-safe.
- **Copyprivate(*list*)**: used to broadcast values obtained by a single thread directly to all instances of the private variables in the other threads.

c/c++:

```
#pragma omp parallel [clauses]
{
    #pragma omp single [clauses]
    ...
}
```

Fortran:

```
!$omp parallel [clauses]
    !$omp single [clauses]
    ...
!$omp end single
!$omp end
parallel
```





Workshare Construct

- Fortran only
- Divides the execution of the enclosed structured block into separate units of work
- Threads of the team share the work
- Each unit is executed only once by one thread

- Allows parallelisation of

- array and scalar assignments
- WHERE statements and constructs
- FORALL statements and constructs
- parallel, atomic, critical constructs

```
!$omp workshare  
...  
!$omp end workshare  
[nowait]
```



```
Program WSex
  use omp_lib
  implicit none

  integer i
  real a(10), b(10), c(10)
  do i=1,10
    a(i)=i
    b(i)=i+1
  enddo

  !$omp parallel shared(a, b, c)
  !$omp workshare
  c=a+b
  !$omp end workshare nowait
  !$omp end parallel

end program WSex
```





References

1. <http://openmp.org>
2. <https://computing.llnl.gov/tutorials/openMP>
3. http://www.openmp.org/mp-documents/OpenMP4.0RC1_final.pdf
4. Michael J. Quinn, Parallel Programming in C with MPI and OpenMP, Mc Graw Hill, 2003.





Thank you!

